

# Dhruv Nirmal

Melbourne, Australia | [dhruvnirmal2111@gmail.com](mailto:dhruvnirmal2111@gmail.com) | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

7 July 2026

Hiring Team, Data Science Chapter, AGL, Australia

Dear Hiring Team,

I'm writing about the Data Scientist role in AGL's Data Science Chapter, partnering with Acquisition Performance and Enablement and Product and Pricing Analytics. What draws me to it is the shape of the work itself: building a model is only half the job, and the role clearly expects the other half too, sitting with commercial stakeholders and turning what the model finds into a decision they can act on. That is exactly how I already work.

At Purchasing Index Data Analytics, part of the Comprara Group, I owned pricing and spend analytics for a restaurant chain operating across Australia and New Zealand as the sole analyst on the account. I built a pricing-decision dashboard that tracked price movement and cost leakage down to the product, venue and supplier level, ran the relationship directly with the Chief Procurement Officer and category managers, and turned a vague sense of overspending into roughly A\$2M a year in identified savings in their largest category. That project sits close to the pricing and commercial decision work this role is built around.

On the modelling side, I led the company's first machine learning classification project, building a one-vs-all classifier over roughly A\$12B of client spend across five years, with proper feature engineering, cross-validation, hyperparameter tuning and class-imbalance handling, cutting a categorisation cycle that used to take an account manager over a month down to a single day's model run. I've also built Prophet forecasting models with external demand drivers for inventory planning, and I use AI tools like Claude in production, not just to experiment, including an agentic reporting workflow now used across our whole analyst team.

AGL sits at the centre of Australia's energy transition, and hiring two data scientists into a chapter that spans acquisition and pricing tells me this is a team backing data to drive real commercial decisions at scale, not just report on them. The toolkit behind classification, forecasting, anomaly detection and pricing analysis on large, messy datasets carries directly across industries, and I'd bring the same habit of framing an ambiguous business question, building a model that actually holds up, and explaining it in plain language to the people who have to act on it.

I'd welcome the chance to discuss how I could contribute to the chapter.

Sincerely,

Dhruv Nirmal

[dhruvnirmal2111@gmail.com](mailto:dhruvnirmal2111@gmail.com) | +61 406 259 619