Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | LinkedIn | GitHub

PROFESSIONAL SUMMARY

I'm a data scientist who builds pricing, forecasting and machine learning models and makes sure the results actually change a commercial decision. I work across large, messy datasets: I built a classification model over roughly **A$12B** of client spend, Prophet forecasting models driven by external demand and price signals, and a pricing analysis that helped a client find around **A$2M** in savings in one category, presented directly to commercial stakeholders up to CFO level. I like owning a modelling problem end to end, framing it, engineering the features, tuning and validating the model, then explaining what it means in plain language to the people who act on it. I'm looking to do exactly that kind of hands-on pricing and analytics work where the model feeds a real business decision.

KEY SKILLS

Machine Learning & Statistical Modelling: scikit-learn, Classification, Logistic Regression, Time-Series Forecasting (Prophet), Anomaly Detection, Feature Engineering, Cross-Validation, Hyperparameter Tuning, Class-Imbalance Handling, Statistical Analysis, Hypothesis Testing

Programming & Big Data: Python, SQL, PySpark, Databricks, Azure, Snowflake, Large & Complex Datasets, ETL

Pricing & Commercial Analytics: Pricing Strategy, Price-Movement & Cost Analysis, Sales & Market Data Analysis, Demand Forecasting, Stakeholder Communication

Visualisation & Storytelling: Tableau (16+ dashboards), Power BI, Data Storytelling, Executive Dashboards

AI & Automation: LLM & agentic workflow automation (Claude, ChatGPT), AI-output QA

EXPERIENCE

Data Scientist, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

A data and analytics consulting firm. I build pricing, forecasting and ML models on large client datasets and work directly with commercial stakeholders, from category managers to CFOs, to turn what the models find into pricing and commercial decisions.

- Owned pricing and spend analytics for a restaurant chain (AU/NZ) as sole analyst: built a pricing-decision dashboard and supplier tracker surfacing price movement and cost leakage down to product, venue and supplier level, helping the client identify roughly **A$2M a year** (about 30%) in savings in their largest category, presented directly to the CFO, Chief Procurement Officer and category managers.
- Led the company's first ML classification project: a one-vs-all binary classifier (TF-IDF, feature engineering, class-imbalance handling, cross-validation, hyperparameter tuning) built over **~A$12B** of client spend across five years, cutting a manual categorisation cycle from 1 to 1.5 months down to a single day's model run; the documented pipeline is now rolled out across the wider client base.
- Built a Prophet time-series forecasting model incorporating external demand and price drivers (freight trends, input pricing) to forecast requirements three months ahead, achieving a 12.5 to 14% error margin comfortably within the client's tolerance, delivered through a dashboard.

- Own the end-to-end data for the company's largest client account across large, complex datasets (5 global regions, 13 sub-regions, ~6M invoice rows refreshed weekly): standardised conventions and automated validation to keep the data reliable, and cut weekly processing from ~4-5 hours to a ~75-minute run.
- Built a company-wide AI automation adopted across the whole analyst team: an LLM-based workflow that auto-generates client reports and saves roughly 12 to 15 hours a week, with a multi-agent QA step that checks outputs before they reach a client, so it stays accurate and governed at scale.
- Built and maintain 20+ automated SQL and Python reporting processes handling millions of transactions monthly, and 16+ Tableau dashboards that stakeholders use to track KPIs and act on cost and pricing decisions.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline in **Databricks with PySpark** on large-scale data, improving detection accuracy about 20% and engineering scalable transformation layers for cloud processing.
- Deployed a Power BI analytics solution adopted by senior Department of Transport stakeholders, presenting model architecture and outcomes to cross-functional, non-technical audiences.

Research Data Analyst, Terminal Ballistics Research Laboratory (TBRL, DRDO)

Jan 2021 to Jul 2021

- Built statistical models to predict blast-wave noise levels, improving prediction accuracy about 20%, and compiled findings into a structured technical report (MATLAB, Excel).

PROJECTS

MyFacit: Analytics and Decision-Support SaaS for Hospitality (Forecasting, Pricing, BI)

- Built a centralised analytics platform for a small hospitality business, integrating 8+ data sources into one consistent view with dashboards, short-term forecasting and insights for pricing, product mix and operations.
- In pilot with a local business as design partner and in active use, with weekly sessions to refine it against real-world data, turning pricing and demand signals into decisions a non-technical operator can act on.

Job Hunt OS: Agentic AI Career System (Python, Claude Code, Multi-Agent)

- Designed and built a multi-agent AI system (orchestrator plus specialist agents) that tailors resumes to a job description, drafts cover letters and runs a two-track outreach process, with a human review gate at every output stage.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Coursework: Statistics I/II, Machine Learning, Applied Forecasting, High Dimensional Analysis, Communicating with Data

Bachelor of Engineering, Thapar University | May 2017 to May 2021