

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I'm a data scientist who builds pricing, forecasting and ML models that feed real commercial decisions, and works with stakeholders to make sure they land. I've built a pricing analysis that helped a client find around **A\$2M** in savings, Prophet forecasting models, and an ML classifier over roughly **A\$12B** of spend. I work mainly in Python and SQL, and have used Databricks and PySpark for large-scale data.

KEY SKILLS

ML & Statistical Modelling: scikit-learn, Classification, Logistic Regression, Time-Series Forecasting (Prophet), Anomaly Detection, Feature Engineering, Cross-Validation, Hyperparameter Tuning, Class-Imbalance Handling

Programming & Big Data: Python, SQL, PySpark, Databricks, Azure, Snowflake, Large & Complex Datasets, ETL

Pricing & Commercial Analytics: Pricing Strategy, Price-Movement & Cost Analysis, Sales & Market Data, Demand Forecasting, Stakeholder Communication

Visualisation: Tableau (16+ dashboards), Power BI, Data Storytelling

AI & Automation: LLM & agentic workflow automation (Claude, ChatGPT), AI-output QA

EXPERIENCE

Data Scientist, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

- Built a pricing-decision dashboard for a restaurant-chain client (AU/NZ), surfacing price movement and cost leakage by product, venue and supplier, and presented to the CFO and category managers; helped identify roughly **A\$2M** (about 30%) in savings in their largest category.
- Led the company's first ML classifier (one-vs-all, TF-IDF, feature engineering, class-imbalance handling, cross-validation, hyperparameter tuning) over **~A\$12B** of spend, cutting a manual categorisation cycle from over a month to a single day's model run.
- Built a Prophet forecasting model using external demand and price drivers to forecast requirements three months ahead, within the client's error tolerance.
- Own large, complex client data (5 regions, 13 sub-regions, ~6M rows refreshed weekly): standardised conventions and automated validation, cutting weekly processing from ~4-5 hours to ~75 minutes.
- Built a company-wide LLM automation adopted across the analyst team that auto-generates client reports (saving ~12-15 hours a week), with a multi-agent QA step checking outputs before delivery.
- Built 20+ automated SQL and Python reporting processes and 16+ Tableau dashboards stakeholders use to track KPIs and inform pricing and cost decisions.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline in **Databricks with PySpark** on large-scale data (about 20% accuracy improvement), and deployed a Power BI solution adopted by senior Department of Transport stakeholders.

PROJECTS

MyFacit: Analytics & Decision-Support SaaS for Hospitality

- Centralised analytics platform integrating 8+ data sources, with dashboards, forecasting and pricing/product-mix insights; in pilot with a local business and in active use.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Bachelor of Engineering, Thapar University | May 2017 to May 2021