

# Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

Hiring Team, Coterie Insurance

Dear Hiring Team,

I'm writing about the Data Analyst role at Coterie. From the job description, this role sits exactly where I like to work: SQL against complex, real datasets, recurring reporting that leaders actually check, and forward-looking modelling layered on top, all in service of a business decision rather than a dashboard for its own sake. I also work fully remote today, and I've built the habits that make that work: clear written documentation, proactive updates, and validation checks that catch problems before a stakeholder ever sees them, so a distributed team never has to chase me for status.

The clearest proof of that is a multi-venue restaurant chain across Australia and New Zealand that I ran as the sole analyst on the account. I segmented their spend by product, venue and supplier to see where prices were moving and where money was leaking out of the category, and that work surfaced close to A\$2 million in savings in one of their largest spend categories over a year, reported directly to their Chief Procurement Officer and category managers in recurring review meetings. It's the same discipline the JD calls portfolio segmentation and partner performance, applied to a different market. I've also standardised how data reaches me for the company's largest client, coordinating 13 regions onto one intake format and building automated checks so bad data gets caught before it reaches a report, which is the same data-integrity bar this role holds.

On the technical side, I built a time-series forecasting model in Prophet for a fresh-produce client, giving their operations team a three-month forward view of inventory needs within a 12.5 to 14% error margin, close enough to plan against. I've also worked hands-on with Databricks and PySpark on a layered lakehouse pipeline during a data engineering internship, building the kind of bronze-to-gold transformation logic this role's Medallion layers describe, and I use Python and Tableau or Power BI daily to turn that kind of modelling into something a non-technical stakeholder can act on in a minute, not a meeting.

What draws me to Coterie specifically is the model itself: taking something that usually takes weeks of underwriting back-and-forth and turning it into a bindable quote in under a minute. That's a data-integrity problem as much as a product one, and it's the exact kind of problem I've spent my career on, taking unfamiliar, messy, multi-source data (procurement invoices, supplier tenders, POS and payroll feeds) and turning it into something a non-specialist can trust and act on fast. Moving that discipline into insurance data feels like a direct next step rather than a stretch.

I'd welcome the chance to talk through how this experience maps onto Coterie's data. Thank you for your time and consideration.

Sincerely,

Dhruv Nirmal