

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I design the SQL data models and pipelines that turn large, messy source data into reliable, reusable datasets other teams can build reporting on. In my current role I own the modelling and validation behind a 6M-row, multi-region client account, build the 20+ automated pipelines and 16+ Tableau and Power BI dashboards that keep that data trustworthy and usable, and coordinate with stakeholders across regions so everyone works from one consistent set of definitions; on a separate account this kind of analysis helped a client find roughly A\$2M in savings. I also design and maintain MyFacit, a live product where I own the data model and weekly metrics end to end, and I work well in distributed teams, documenting decisions clearly so pipelines and definitions hold up across time zones.

KEY SKILLS

SQL & Data Modelling: SQL (SQL Server, Snowflake), modular and reusable SQL transformations, data modelling, query optimisation, Python

ETL & Data Pipelines: automated Python/SQL pipelines, API-based ingestion, batch scheduling, full and incremental refresh, data validation, standardised intake

Data Warehousing & Cloud: Snowflake, AWS (S3, EC2, IAM), Azure, Terraform (IaC), relational database design, Docker

BI & Reporting: Tableau, Power BI, KPI design and monitoring, dashboard development, self-service reporting

Tools & Practices: Git, documentation, stakeholder collaboration, data-quality/validation processes

EXPERIENCE

Data & Analytics Engineer, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia (remote-capable client analytics consulting)

- Design and maintain the SQL data models and transformation logic behind the firm's largest client account, turning raw invoice, PO and lookup data from 5 global regions (about 6M rows) into one standardised, reusable dataset for downstream reporting.
- Build and maintain 20+ automated ETL pipelines (Python, SQL) across 5 enterprise clients, processing millions of procurement transactions monthly on scheduled, Azure-hosted infrastructure.
- Own data quality on that largest account: automated validation (row and hour counts, data-type and column-consistency checks) that catches discrepancies before delivery, replacing a manual, error-prone upload process.
- Build and maintain 16+ Tableau and Power BI dashboards that give procurement and finance stakeholders self-service KPI reporting instead of ad hoc requests.
- Coordinate with stakeholders and analysts across 13 sub-regions to standardise file submission, naming conventions and metric definitions so every region feeds one consistent model.
- Leading the automation behind a company-wide, self-service pipeline-refresh app across 26 client accounts (Python, SQL, Git), turning a manual, specialist refresh into a one-click run with separate controls for ingestion, validation and processing.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline (Databricks, PySpark) and engineered scalable transformation layers for government procurement data, improving detection accuracy by about 20%.
- Delivered a Power BI analytics solution adopted by senior Department of Transport stakeholders, presenting the architecture and outcomes to cross-functional stakeholders.

Research Data Analyst, Terminal Ballistics Research Laboratory (TBRL), DRDO

Jan 2021 to Jul 2021

- Applied statistical modelling and signal filtering to blast-wave data to predict noise levels, improving prediction accuracy by about 20%, and compiled findings into a 70-page technical report.

PROJECTS

Cloud Data Platform (Terraform, AWS, Snowflake)

- Provisioned cloud infrastructure as code (S3, IAM, EC2) with Terraform and built an automated ingestion pipeline from external APIs into a Snowflake data warehouse.
- Modular, multi-environment design (dev, staging, prod), hands-on with core data warehousing and infrastructure-as-code concepts.

MyFacit, profitability analytics SaaS for independent cafes (personal build)

- Own the data model and metric definitions end to end for a live multi-tenant product that unifies POS transactions, staff wages and 14 suppliers' invoices into one weekly profitability view.
- Define and maintain the metrics it reports (revenue, wage cost %, food cost %, gross margin) plus revenue forecasting and supplier price-creep detection, live in day-to-day use with a design-partner cafe (Neighbours Cafe, St Kilda).

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Coursework: Statistics I & II, Machine Learning, Communicating with Data, Applied Forecasting

Bachelor of Engineering, Thapar University | May 2017 to May 2021