

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

Hiring Team, Deel

Dear Hiring Team,

I'm writing about the Analytics Engineer role at Deel. Most of my day to day is exactly this work: taking raw, messy source data and turning it into SQL models and pipelines that other people can trust and build reporting on, then keeping those pipelines and definitions consistent as the data keeps growing. A global, remote-first platform like Deel, where the same underlying data has to stay accurate and consistent across a huge number of countries and systems, is the kind of scale where that discipline actually matters, and it's the part of the job I enjoy most.

The clearest proof of that is the largest client account I run in my current role at Purchasing Index Data Analytics. It's a multi-region account spanning 5 regions and 13 sub-regions, close to 6 million invoice rows feeding in every week, each region originally submitting data in its own format through a manual upload process. I designed the SQL data model and transformation logic that turns all of that into one standardised, reusable dataset, replaced the manual upload with automated ingestion, and built the validation checks that catch discrepancies before they ever reach a client. Getting 13 different regions to feed one consistent model, with the same definitions and the same quality bar, is the same problem Deel's analytics engineers solve at an even bigger scale, and it's the piece of my experience I'd bring first.

Alongside that, I've built and maintain 20+ automated Python and SQL pipelines and 16+ Tableau and Power BI dashboards across five enterprise clients, and I provisioned a Snowflake data warehouse as infrastructure as code with Terraform on a personal build, so the warehousing side of this role isn't new to me either. I also design and own the data model for MyFacit, a live product I built that unifies point-of-sale, payroll and supplier data for independent cafes into one weekly profitability view. I decide what it tracks and define its metrics myself, the same instinct for keeping one dataset and one set of definitions consistent that I'd bring to Deel's models.

What draws me to Deel specifically is the combination of the mission and the way the team works: payroll and compliance across 100+ countries is fundamentally a data-consistency problem at scale, solved by a fully distributed team, which is exactly the environment I already work well in. I write things down, keep decisions and pipeline logic documented, and give proactive updates by default, because that's what makes distributed collaboration actually work, and it's how I'd plug into a team spread across as many time zones as Deel's.

I'd welcome the chance to talk through how this experience applies to Deel's data models. Thank you for your time and consideration.

Sincerely,

Dhruv Nirmal