

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

Hiring Team, GitHub | Data Analyst, Trust & Safety (Copilot Abuse)

I am writing about the Data Analyst role on the Trust and Safety team focused on Copilot abuse. The work described, turning product telemetry and account signals into abuse-pattern intelligence that drives enforcement and policy decisions, maps precisely to how I work every day: I dig into large operational datasets to separate normal activity from problematic patterns, build the dashboards and automated workflows that surface those findings to the people who act on them, and evaluate AI-generated outputs critically before they reach stakeholders. The cross-functional consulting model GitHub describes, partnering with Security, Legal, and Product to align on data definitions and translate findings into decisions, is the mode I operate in as a client-facing analyst, not a stretch I would be adjusting to.

The strongest proof I can offer is anomaly detection on real operational data. At Purchasing Index, I applied anomaly detection to enterprise procurement transaction datasets to surface cost-leakage and irregular spending patterns, contributing to roughly 15% savings for clients by distinguishing normal supplier behaviour from genuinely problematic activity. I own the full pipeline for the largest client account (roughly 6M invoice rows across 13 global sub-regions), including automated data-quality controls that flag integrity issues before they reach any stakeholder. At the Victorian Centre for Data Insights, I built a distributed anomaly-detection pipeline in Databricks and PySpark on government procurement data, improving detection accuracy by roughly 20%. Both projects are the same analytical pattern as Copilot abuse detection: define what normal looks like, build the detection logic, and surface the deviations clearly. The client relationship I ran independently for a restaurant chain across AU and NZ produced roughly A\$2M in identified savings over a year by surfacing cost-leakage to the product and venue level, and I ran all stakeholder communication directly with the CFO, Chief Procurement Officer, and category managers.

The second signal I want to put on the table is hands-on AI and LLM system building. I have designed and shipped production LLM and multi-agent systems: a call-transcript analyser that scores quality, expectation fulfilment, and sentiment across client engagements and routes findings to stakeholders automatically (placed 1st at an internal AI hackathon, then adopted into the company workflow), a 5-agent plus orchestrator data QA pipeline that validates client datasets before delivery, and agentic report-generation workflows used by the entire analyst team. This means I understand how AI tools behave from the inside: how to evaluate their outputs for accuracy and consistency, where they can be gamed or misused, and how to build reliable detection logic around AI-generated content. That perspective is directly relevant to reasoning about Copilot prompt, completion, and usage-pattern abuse in ways that a pure data analyst background would not be.

GitHub and Copilot represent exactly the kind of problem I find most compelling: large-scale platform data, a fast-evolving set of abuse patterns, and real downstream impact on product integrity and user trust. My background is procurement analytics rather than developer-platform Trust and Safety, but the analytical methods are the same and the AI-product fluency is genuine. The combination of behavioural-pattern detection on real operational data, production LLM building, and consulting delivery across senior stakeholders is the specific mix this role calls for, and I would welcome the chance to discuss it.

Regards,

Dhruv Nirmal