

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

As a client-facing data analyst, I turn complex datasets into ROI stories that land with executive audiences. At a data and analytics consulting firm, I run five enterprise client accounts solo, owning each relationship from an ambiguous initial brief through to structured analysis, dashboard delivery, and direct presentation to C-suite stakeholders including CFOs and Chief Procurement Officers. On one account, translating a restaurant chain's vague sense of overspending into roughly **A\$2M a year** in identified savings, by surfacing price movement and cost leakage to the product and venue level, captures exactly the kind of work this role is built around. I work daily with SQL and Python over large, messy datasets, build executive-facing dashboards in Tableau and Power BI, and integrate AI-assisted tooling into analytics workflows as a genuine productivity lever. The Data Analyst role on HTB's Federal Sales team combines the two things I do best: owning client relationships at the executive level, and using data to tell a clear ROI story.

KEY SKILLS

Programming and Analytics: Python, SQL, Snowflake, Statistical Analysis, Anomaly Detection, Forecasting, Feature Engineering, Excel (Pivot Tables, VLOOKUP, ad-hoc modeling)

Visualisation and BI: Tableau (16+ dashboards), Power BI, Executive Dashboards, Data Storytelling, KPI Analytics

AI and LLM Tooling: Claude, ChatGPT, Agentic AI Workflows, Multi-Agent Orchestration, LLM-Accelerated Analysis, NLP/Sentiment Analysis, AI Output Evaluation

ML and Data Science: scikit-learn, TF-IDF/NLP, Logistic Regression, Prophet Forecasting, Classification, Cross-Validation

Cloud and Data Engineering: Azure, AWS (S3, EC2, IAM), Databricks, ETL Pipeline Design, Data Quality/Validation, KNIME

Consulting: Client-Facing Delivery, Stakeholder Management, ROI Framing, Executive Presentations, Problem Framing

EXPERIENCE

Data Analyst, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

A data and analytics consulting firm serving enterprise clients across Australia and New Zealand.

Independently own and run 5 enterprise client accounts as the sole analyst on each: scope ambiguous problems, gather data and context from client teams, run the analysis, build the dashboards, and present recommendations directly to C-suite stakeholders (CFOs, Chief Procurement Officers). Manage delivery timelines, expectations, and a regular meeting cadence across all five relationships.

- Owned the end-to-end analytics account for a restaurant chain (AU/NZ) as sole analyst: ran bi-weekly progress meetings with CFO, CPO, category managers and general managers; structured ambiguous cost problems into testable hypotheses; surfaced **~A\$2M/year** in savings (~30%) in their largest category through price-movement tracking and cost-leakage analysis delivered via interactive dashboards.
- Built and maintained 16+ Tableau dashboards for enterprise clients, translating procurement and spend data into executive-level views and operational drill-downs that inform supplier negotiations and category decisions.
- Automated client reporting workflows (Python, SQL) and integrated LLM-assisted tooling (Claude, ChatGPT) to accelerate analytics script development; built a company-wide agentic report-generation workflow that saves **~12-15 hours per week** across ~25-30 client accounts by replacing manual per-client report assembly.

- Performed anomaly detection on procurement datasets (Python, SQL) to identify cost-leakage patterns and surface high-value irregularities as actionable findings for client teams.
- Led the company's first ML spend-categorisation project: designed a one-vs-all binary classifier (TF-IDF, Python, scikit-learn) over **~A\$12B of client spend** across 5 years, cutting the categorisation cycle from **1-1.5 months to a single day's model run**; pipeline now rolled out across clients.
- Built a Prophet time-series forecasting model (Python) incorporating external price drivers (sea-freight trends) to predict 3-month raw material requirements for a fresh-produce client, achieving a **12.5-14% error margin** within the stated tolerance.
- Own the end-to-end data pipeline for the company's highest-billing client: spend data across 5 global regions and 13 sub-regions (~6M invoice rows, refreshed weekly); replaced manual upload with automated email-based ingestion (Microsoft Graph API, SQL Server), cutting weekly processing from ~4-5 hours to a ~75-minute automated run.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline (Databricks, PySpark) improving detection accuracy 20%; engineered scalable transformation layers for cloud procurement data.
- Deployed a Power BI analytics solution adopted by senior Department of Transport stakeholders; presented model architecture and outcomes to cross-functional non-technical audiences.

Research Data Analyst, Terminal Ballistics Research Laboratory (TBRL), DRDO

Jan 2021 to Jul 2021

- Analysed blast-wave sensor data using statistical models (MATLAB, Excel) and compiled findings into a 70-page structured report, translating quantitative outputs into clear recommendations for defence engineering teams.

PROJECTS

LLM Call-Transcript Analyser (Python, LLM Workflows, Sentiment Analysis, Agent Orchestration)

- Built for an internal AI hackathon (1st place), then adopted into the company's workflow: a pipeline that ingests call transcripts, scores each engagement on call quality, expectation alignment, and sentiment, then routes findings to relevant stakeholders via automated notifications and an engagement log.
- Demonstrates behavioral and engagement analysis in production: scoring completion signals, identifying sentiment patterns, and surfacing findings to decision-makers automatically.

MyFacit: Analytics and Decision-Support SaaS (Data Integration, Forecasting, BI, AI Automation)

- Designed and built a centralised analytics platform for a hospitality business, integrating 8+ data sources and AI-driven workflows to automate reporting and surface actionable insights for pricing, product mix, and operational planning; in active pilot with a real client, with weekly feedback loops.
- Built the data governance and quality layer: standardised definitions and validation checks that reconcile 8+ disparate sources into one consistent, trustworthy model for decision-making.

Cloud Data Platform with Terraform (AWS, Snowflake, Terraform, Docker, Power BI)

- Provisioned full cloud infrastructure as code (Terraform: S3, IAM, EC2); built an automated ingestion pipeline from external APIs to Snowflake; modular multi-environment design (dev/staging/prod) with Power BI reporting layer on top.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Coursework: Statistics I/II, Machine Learning, Applied Forecasting, High Dimensional Analysis, Communicating with Data

Bachelor of Engineering, Thapar University | May 2017 to May 2021