

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I build and maintain the pipelines that pull data from different source systems into one reliable base, then turn it into the dashboards and reporting that business teams use to make decisions. In my current role at Purchasing Index Data Analytics I own the end-to-end pipeline for the firm's largest client, close to six million transaction rows drawn from five global regions, and I have built and maintained 20+ automated ETL pipelines across five enterprise clients, backed by Tableau reporting for stakeholders. On one engagement my analysis and dashboards helped a client identify close to A\$2 million in savings. I also design and run MyFacit, a live product I built that unifies point-of-sale, payroll and supplier invoice data for independent cafes into a single weekly profitability view, direct experience integrating retail and operational systems into one source of truth. I want to bring that pipeline-to-insight discipline to the Data & Analytics Engineer role at Johnson Fitness & Wellness.

KEY SKILLS

Data Engineering & Pipelines: Python, SQL, ETL/ELT pipeline design, API integration (Microsoft Graph API), batch automation, full and incremental refresh, data modelling, data validation and quality monitoring, documentation

Cloud & Data Warehousing: Snowflake, Databricks, PySpark, AWS (S3, EC2, IAM, Lambda), Azure, Terraform (IaC), Docker

BI & Reporting: Tableau, Power BI, KPI reporting, dashboard design, data storytelling

Analytics & Forecasting: Prophet time-series forecasting, anomaly detection, statistical analysis, scikit-learn

Tools & Practices: Git / version control, KNIME, stakeholder communication

EXPERIENCE

Data & Analytics Engineer, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

A data and analytics consultancy building pipelines and reporting across supplier, transaction and procurement systems for enterprise clients in Australia and New Zealand. Selected engagements:

Multi-source pipeline integration: largest enterprise client (multi-region)

- Problem: the account ran on a manual upload-and-check process, about 4 to 5 hours every week, across close to six million invoice rows spanning 5 global regions and 13 sub-regions, each submitting files in its own format.
- Approach: replaced it with automated email-based ingestion (Microsoft Graph API into a remote SQL Server via batch scripts), standardised file submission and naming conventions so every region feeds one consistent intake, and built automated validation (row and hour counts, data-type and column-consistency checks) that drops stray columns and logs discrepancies for review.
- Result: cut weekly processing from 4 to 5 hours down to about 75 minutes, roughly 3 to 3.5 hours saved every week, with one standardised pipeline now feeding downstream reporting for the company's highest-billing client.

Pricing and cost analytics: multi-venue hospitality client (~A\$2M savings)

- Problem: the client suspected significant supplier overspend but had no dataset or visibility into where cost was leaking across products, venues and suppliers.
- Approach: sole analyst on the account; built a roughly 3,000-item product catalogue from scratch, segmented spend by product, venue and supplier, and built a Tableau dashboard tracking price movement and cost leakage, presenting findings directly to the client's Chief Procurement Officer and category managers.
- Result: helped the client identify close to A\$2 million (about 30%) in savings in their largest category over a year, used directly in supplier negotiations.

Enterprise ETL pipelines across 5 clients

- Problem: each client's procurement and transaction data needed repeatable, reliable ingestion from separate source systems.
- Approach: built and maintain 20+ automated ETL pipelines in Python and SQL, deployed on Azure-hosted servers with scheduled automation.

- Result: cut manual reporting effort by about 40% and improved turnaround by about 50% across all 5 accounts.

Self-service pipeline refresh app (26 client accounts, in progress)

- Problem: refreshing each of the firm's 26 client datasets was a manual, specialist task that tied up the analytics team.
- Approach: leading the automation behind a company-wide app that turns each client refresh into a one-click run, with separate controls for ingestion, validation and downstream processing, both full and incremental, version-controlled in Git with maintained documentation and clean underlying databases.
- Result: on track to save about 60 hours a month across the team once complete.

Reporting, BI and data quality: developed and maintain 16+ Tableau dashboards for procurement and finance stakeholders, translating transaction and supplier data into KPIs and cost-optimisation opportunities; ran anomaly detection across procurement datasets to flag cost-leakage patterns, contributing to about 15% in client savings.

Forecasting for inventory planning: built a Prophet time-series forecasting model for a fresh-produce client, projecting raw-material and chemical inventory needs about three months ahead and incorporating external price drivers such as sea-freight trends, delivered on a dashboard within a 12.5 to 14% error margin, well inside the client's tolerance.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline in Databricks using PySpark, engineering scalable transformation layers for government procurement data and improving detection accuracy by about 20%.
- Delivered a Power BI analytics solution adopted by senior Department of Transport stakeholders, presenting the architecture and outcomes to cross-functional stakeholders.

PROJECTS

MyFacit: multi-tenant profitability SaaS for independent cafes (personal build)

- Built a product that unifies POS transactions, staff wages and 14 suppliers' invoices into one weekly profitability view: revenue, wage cost %, food cost %, gross margin and a recommended action for the week ahead.
- Includes revenue forecasting, menu engineering and supplier price-creep detection; live with a St Kilda cafe as the design partner, with weekly meetings to refine it against real, messy operational data. Direct experience with the kind of multi-source, retail-adjacent data integration this role needs.

Cloud Data Platform with Terraform (personal build)

- Provisioned cloud infrastructure as code with Terraform (AWS S3, IAM, EC2, networking) and built an automated ingestion pipeline from external APIs into Snowflake, with Docker for a modular, multi-environment setup (dev, staging, prod).

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Bachelor of Engineering, Thapar University | May 2017 to May 2021