

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I like taking an ambiguous business question all the way through: exploring the data, engineering the right features, and shipping a working Python model rather than stopping at a proof of concept. At PI Data Analytics, a consulting firm, I built a scikit-learn classification pipeline that brought structure to roughly A\$12B of a client's spend history, and have since built Prophet demand forecasting and internal generative AI and agentic workflows now in everyday use across the business. I moved into data science through a Master of Data Science at Monash after an engineering background, and keep that curiosity active with a self-directed DS study routine and an agentic system I built to run my own job search. I'd like to bring that hands-on, curious approach to Keypath's Strategy & Insights team.

KEY SKILLS

Machine Learning & Statistics: Python (pandas, NumPy, scikit-learn), classification (TF-IDF, logistic regression), time-series forecasting (Prophet), feature engineering, cross-validation, hyperparameter tuning, class-imbalance handling, anomaly detection

Generative AI & LLMs: LLM-powered agents (Claude, ChatGPT), prompt engineering, multi-agent orchestration, agentic workflow design, critical evaluation of AI/LLM output

Data, Cloud & Visualisation: SQL, EDA, data wrangling, Snowflake, Git/GitHub, Azure (light exposure), Databricks/PySpark, Power BI, Tableau, presenting to non-technical stakeholders

EXPERIENCE

Data Analyst, PI Data Analytics (Comprara Group) | Jun 2024 to Present | Melbourne

A data and analytics consulting firm serving enterprise clients across Australia and New Zealand.

- Led the firm's first ML classification project: a one-vs-all classifier (scikit-learn, TF-IDF, feature engineering, class-imbalance handling, cross-validation, hyperparameter tuning) over roughly A\$12B of a client's spend across five years, cutting a manual categorisation cycle from about six weeks to a single day's model run; documented and rolled out across the wider client base.
- Built a Prophet time-series demand-forecasting model for a fresh-produce client, incorporating external drivers such as sea-freight trends to forecast inventory 3 months ahead within a 12.5 to 14% error margin, delivered through a dashboard.
- Designed internal generative AI and agentic workflows: an agentic report-generation system (Claude) rolled out across roughly 25 to 30 accounts and saving the analyst team about 12 to 15 hours a week, a multi-agent data-QA workflow, and an LLM call-transcript sentiment analyser that placed first at an internal AI hackathon and was adopted into the team's workflow.
- Own pricing analysis and dashboards end to end for a multi-venue restaurant chain (AU/NZ) as sole analyst, tracking price movement and cost leakage to product, venue and supplier level and presenting in plain language to senior stakeholders including the Chief Procurement Officer; helped identify close to 30% (about A\$2M) in savings in one of their largest spend categories over a year.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI) | Aug 2023 to Nov 2023 | Melbourne

- Built a distributed anomaly-detection pipeline (Databricks, PySpark) on government procurement data, improving detection accuracy by about 20%; deployed a Power BI solution adopted by senior Department of Transport stakeholders.

PROJECTS

- **Job Hunt OS: agentic AI career system (Python, Claude Code, multi-agent orchestration):** a file-based, multi-agent system (orchestrator plus specialist agents) that tailors resumes to a JD and runs a structured job search with a human review gate, on GitHub. It also runs a self-directed DS curriculum I built in (daily study logs, weekly mock interviews), how I develop my DS skills outside client work.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Bachelor of Engineering, Thapar University | May 2017 to May 2021