

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

Hiring Team, Strategy & Insights, Keypath Education, Melbourne

Dear Hiring Team,

I'm writing about the Associate Data Scientist role on the Strategy & Insights team. I came to data science by way of an engineering degree and a Master of Data Science at Monash, then spent the last two years putting that training to work on real commercial problems rather than only on coursework, which is exactly the mix this role is built for. I like taking a messy, ambiguous question, working through the data with pandas and scikit-learn, and landing on a model that a business can actually use, and I'd like to keep building that muscle inside a team that treats data science as core to how the company operates rather than a side function.

The clearest proof of that is a classification pipeline I led at PI Data Analytics, a data and analytics consulting firm, over roughly A\$12B of a client's spend history. Only about 60 to 65% of that spend was reliably categorised when I started. I built a one-vs-all classifier with proper feature engineering, class-imbalance handling, cross-validation and hyperparameter tuning, and paired it with a Pareto approach so the model handles the long tail while a person reviews the highest-value items. It turned a categorisation cycle that used to take an account manager a month to a month and a half into a single day's model run, and it's since been documented and rolled out across other client accounts. I've applied the same discipline, exploratory analysis, feature engineering, honest validation, to a Prophet-based demand forecasting model for a fresh-produce client, forecasting raw-material needs three months out within a tolerance the client was comfortable with.

The other half of what I'd bring is genuine, hands-on experience with generative AI in production, not just using it. I designed an agentic workflow that now auto-generates client progress reports company-wide, saving the analyst team real hours every week, and a multi-agent QA system that checks outputs before they reach a client. Outside work, I built a Job Hunt OS: a multi-agent, Claude-native system, version-controlled on GitHub, that tailors resumes and manages a structured job search end to end. It's also where I run a self-directed data science study routine, daily lessons and weekly mock interviews, which is the clearest evidence I can offer that I keep actively engaging with this field beyond what any one job asks of me.

What draws me to Keypath specifically is the chance to apply that same modelling and GenAI experience to a mission I find genuinely worth working on: helping universities reach and support more students through better-run online programs. My path into data science through a Master's alongside real commercial delivery is exactly the career-transitioner profile this role describes, and I'd welcome the chance to talk through how I can contribute to the team from day one.

Thank you for your time and consideration. I'd welcome the opportunity to discuss this further.

Kind regards,

Dhruv Nirmal