

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I'm a hands-on analyst who'd rather write my own SQL and Python than wait on a pre-built dashboard or someone else's backlog. At PI Data Analytics, a data and analytics consulting firm, I independently manage the full analytics loop for enterprise clients: pulling and joining data from multiple systems myself, building the pipeline, running the statistical work, and distilling it into a recommendation a non-technical stakeholder can actually use. That approach has produced real outcomes, including roughly A\$2M in surfaced savings for a multi-venue restaurant client by tracking price movement and cost leakage down to the product, venue, and supplier level, and a classification pipeline that cut a six-week manual categorisation cycle to a single automated run across a roughly A\$12B spend dataset. I'm looking to bring that same source-to-recommendation ownership to Puffy's experimentation and CRO work.

KEY SKILLS

Analytics & Statistics: SQL (CTEs, window functions, complex joins), Python (pandas, numpy, scikit-learn), hypothesis-driven classification, cross-validation, class-imbalance handling, segment/venue/supplier-level analysis, time-series forecasting (Prophet)

Data Pipelines & Warehousing: Multi-source data integration (APIs, email ingestion, SQL Server, flat files) without engineering support, ETL pipeline design, Snowflake, Azure, AWS (S3, Athena, Glue), BigQuery/GCP (familiar)

AI & Automation: Agentic AI and LLM workflows (Claude, ChatGPT) built on SQL and Python foundations, multi-agent orchestration, prompt engineering, AI-output QA

BI & Stakeholder Communication: Tableau, Power BI, KPI dashboards, client-facing reporting to non-technical senior stakeholders, clear written documentation for distributed teams

EXPERIENCE

Data Analyst, PI Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

Independently manage analytics engagements across 5 enterprise clients for this data and analytics consulting firm (Australia and New Zealand), owning each account's SQL/Python analysis, data pipeline, and stakeholder reporting end to end.

Pricing analytics, multi-venue restaurant group (AU/NZ)

- Problem: the client's category spend was leaking margin across roughly 3,000 products and multiple venues with no visibility into where.
- Approach: owned the account as the sole analyst, built the product catalog from scratch, then a pricing-decision dashboard tracking price movement at the product, venue, and supplier level using direct SQL/Python analysis, and presented findings in recurring meetings with the Chief Procurement Officer and category managers.
- Result: surfaced roughly A\$2M (about 30%) in identifiable savings in the client's largest spend category over a year.

Spend classification and Pareto-driven pipeline, enterprise client (approximately A\$12B spend, 5 years)

- Problem: only about 60 to 65% of a five-year, roughly A\$12B spend dataset was reliably categorised.
- Approach: led the build of a one-vs-all classification pipeline (scikit-learn, TF-IDF plus engineered features on spend value and line count), tuned with cross-validation and explicit class-imbalance handling, paired with a Pareto approach where I manually reviewed the high-value head of spend and let the model classify the long tail.
- Result: cut a categorisation cycle from roughly six weeks of manual account-manager work to about a single day's model run; the pipeline is now the standard being rolled out across other client accounts.

Multi-source pipeline ownership, largest client account (global, 5 regions / 13 sub-regions)

- Problem: the firm's highest-billing client account ran on a manual upload process, costing roughly 4 to 5 hours of manual checking every week with no single connected pipeline.

- Approach: pulled and joined the data myself without waiting on a separate engineering team, replacing the manual process with automated email-based ingestion (Microsoft Graph API into a remote SQL Server batch process), standardised the intake format across all 13 sub-regions, and built automated validation to catch stray columns and discrepancies before they reached the client.
- Result: cut weekly processing from 4 to 5 hours down to about 75 minutes (roughly 3 to 3.5 hours saved weekly) on the firm's highest-billing account.

Internal (own firm, not client-facing): agentic AI and automation

- Built and orchestrate agentic AI/LLM workflows (Claude, ChatGPT) that speed up analytics-script development, run automated multi-agent QA validation on datasets before client delivery, and auto-generate client progress reports, saving the wider analyst team roughly 12 to 15 hours weekly across the firm's client accounts.
- Built a call-transcript sentiment and quality analyzer that took 1st place in an internal AI hackathon and was later adopted into the company's workflow.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline (Databricks, PySpark), improving detection accuracy by roughly 20%.
- Deployed a Power BI analytics solution adopted by senior Department of Transport stakeholders and presented findings to cross-functional audiences.

PROJECTS

Inventory-demand forecasting, fresh-produce client (Python, Prophet)

- Client engagement (PI Data Analytics): built a Prophet time-series model forecasting raw-material and chemical inventory roughly 3 months ahead, engineering external price-driver features (including sea-freight trends) for internationally sourced chemicals, delivered on a client dashboard.
- Delivery/impact: forecast landed within a 12.5 to 14% error margin, comfortably inside the client's tolerance, giving them a real forward planning window.

Facit (myfacit.com), profitability SaaS for independent cafes (personal, forecasting + BI)

- Personal build, not client/employer work: a multi-tenant product giving cafe owners a weekly profitability view, unifying POS, wages, and supplier invoice data into a single "pulse" (revenue, wage cost %, food cost %, gross margin) with one recommended action.
- Includes revenue forecasting and supplier price-creep detection, designed around a real, non-technical operator (Neighbours Cafe, Melbourne, as design partner) and messy real-world data.
- Delivery/impact: live and in active weekly use with the design-partner cafe; turns unintegrated operational data into a Monday-morning decision.

Job Hunt OS, agentic career-search system (personal, Claude Code, Python)

- Personal build: a file-based, multi-agent system (orchestrator plus resume, outreach, and career-coach specialist agents) that tailors resumes to a JD, drafts cover letters, and runs a two-track networking CRM, all without paid APIs.
- Delivery/impact: demonstrates end-to-end agentic AI orchestration and system design over a plain-file data model, applied to a real, ongoing job search.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Statistics I/II, Machine Learning, Communicating with Data, Applied Forecasting, High Dimensional Analysis, Supply Chain Management, Web App Development.

Bachelor of Engineering, Thapar University | May 2017 to May 2021

Statistics, C, C++, Analog & Digital Electronics, Mathematics I/II/III.