

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I'm a hands-on analyst who writes my own SQL and Python rather than waiting on a pre-built dashboard or an engineering queue. At PI Data Analytics, a consulting firm, I own the full loop for enterprise clients: pulling and joining data across systems, running the statistical work, and turning it into a recommendation a non-technical stakeholder can act on, including work that surfaced roughly A\$2M in savings for one client. I want to bring that source-to-recommendation ownership to Puffy's experimentation and CRO program.

KEY SKILLS

Analytics & Statistics: SQL (CTEs, window functions, complex joins), Python (pandas, numpy, scikit-learn), classification, cross-validation, class-imbalance handling, segment/venue/supplier-level analysis, time-series forecasting (Prophet)

Pipelines & Cloud: Multi-source integration (APIs, SQL Server, flat files) without engineering support, ETL design, Snowflake, Azure, AWS (S3, Athena, Glue), BigQuery/GCP (familiar)

AI & Communication: Agentic AI and LLM workflows (Claude, ChatGPT), prompt engineering; Tableau, Power BI; client-facing reporting to non-technical senior stakeholders

EXPERIENCE

Data Analyst, PI Data Analytics (Comprara Group) | Jun 2024 to Present | Melbourne

Independently manage analytics engagements across 5 enterprise clients (AU/NZ) for this consulting firm, owning each account's SQL/Python analysis, pipeline, and stakeholder reporting end to end.

- Surfaced roughly A\$2M (about 30%) in savings in a multi-venue restaurant client's largest spend category, tracking price movement and cost leakage to the product, venue, and supplier level in SQL/Python and presenting to the Chief Procurement Officer and category managers.
- Built a one-vs-all classification pipeline (scikit-learn, TF-IDF plus engineered spend and line-count features, cross-validation, class-imbalance handling) that cut a roughly six-week manual categorisation cycle to a single model run across a A\$12B, five-year spend dataset; now the standard rolling out across other accounts.
- Replaced a manual upload process on the firm's highest-billing client (5 regions, 13 sub-regions) with automated email ingestion (Microsoft Graph API into SQL Server) and validation, cutting weekly processing from 4 to 5 hours to about 75 minutes.
- Internal (own firm): build and orchestrate agentic AI/LLM workflows (Claude, ChatGPT) that speed analytics-script development and run automated QA on datasets before client delivery, saving the analyst team roughly 12 to 15 hours weekly; a call-transcript analyzer I built took 1st in an internal AI hackathon.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI) | Aug 2023 to Nov 2023 | Melbourne

- Built a distributed anomaly-detection pipeline (Databricks, PySpark), improving detection accuracy by roughly 20%; deployed a Power BI solution adopted by senior Department of Transport stakeholders.

PROJECTS

- **Inventory-demand forecasting, fresh-produce client (Python, Prophet):** 3-month time-series forecast with engineered external price-driver features, delivered on a client dashboard within a 12.5 to 14% error margin.
- **Facit (myfacit.com), profitability SaaS for cafes (personal):** unifies POS, wages, and supplier invoices into a weekly profit "pulse" with revenue forecasting and supplier price-creep detection; live with a design-partner cafe.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Bachelor of Engineering, Thapar University | May 2017 to May 2021