

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I'm a data scientist and analytics consultant who runs five enterprise client accounts solo, taking each from a loosely defined problem through to a recommendation the client can act on. I'm at my best where the question isn't clear yet: working out what's actually being asked, tracking down the data and the people who hold the answers, and explaining what it means in plain terms to executives without a technical background. On one account, a restaurant chain's hunch about overspending turned into roughly **A\$2M** a year in savings once I framed the numbers properly. I use AI heavily day to day, but with a critical eye for where it falls short, on a solid Python, SQL and statistics base that includes ML work across about **A\$12B** of client spend. I'm drawn to Quantum for the mix of hands-on consulting and data science, and the chance to bring it into financial services.

KEY SKILLS

Programming and Analytics: Python, SQL, Statistical Analysis, Anomaly Detection, Hypothesis Testing, Feature Engineering, Time-Series Forecasting

AI and LLM: Claude, ChatGPT, Agentic AI Workflows, Multi-Agent Orchestration, AI Output Evaluation, NLP/Sentiment Analysis, Prompt Engineering

Visualisation and BI: Tableau (16+ dashboards), Power BI, Data Storytelling, KPI Analytics, Executive Dashboards

ML and Data Science: scikit-learn, TF-IDF/NLP, Logistic Regression, Prophet, Classification, Cross-Validation, Class-Imbalance Handling

Consulting and Cloud: Client-Facing Delivery, Problem Framing, Stakeholder Communication, QA Documentation, Azure, AWS, Databricks

EXPERIENCE

Data Scientist, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

A data & analytics consulting firm serving enterprise clients across Australia and New Zealand.

Independently own and run 5 enterprise client accounts as the single point of contact: scoping ambiguous problems before there's a defined brief, gathering data and context from client teams, agreeing on the approach, running the analysis, and presenting recommendations directly to C-suite stakeholders (CFOs, CPOs). Manage delivery timelines, expectations, and a regular meeting cadence across all five relationships, and look for ways to extend value within each engagement.

- Owned the end-to-end analytics account for a restaurant chain (AU/NZ) as sole analyst: ran bi-weekly progress meetings with CFO, CPO, category managers and general managers, structured ambiguous cost problems into testable hypotheses, and surfaced **~A\$2M/year** in savings (~30%) in their largest procurement category through price-movement and cost-leakage analysis.
- Designed and deployed a company-wide agentic report-generation workflow (Claude): rolled out across the full analyst team (~25-30 client accounts), replacing manual per-client report assembly and saving **~12-15 hours per week** across the business.
- Led the company's first ML spend-categorisation project: designed a one-vs-all binary classifier (TF-IDF, feature engineering, class-imbalance handling, cross-validation, hyperparameter tuning) over **~A\$12B of client spend** across 5 years, cutting the categorisation cycle from **1-1.5 months to a single day's model run**; the documented pipeline is now being rolled out across clients.
- Developed a multi-agent data QA validation system (5 specialist agents + orchestrator, Claude) to validate client datasets prior to delivery, applying AI-output evaluation at each step to flag where model outputs required human scrutiny before sign-off.
- Built a Prophet time-series forecasting model (Python) for a fresh-produce client, incorporating sea-freight trends and external price drivers to predict 3-month raw material and chemical requirements, achieving a **12.5-14% error margin** comfortably below the client's stated tolerance.

- Own the end-to-end data pipeline for the company's largest, highest-billing client: spend data across 5 global regions and 13 sub-regions (~6M invoice rows, refreshed weekly). Established **data governance** across all 13 regions, standardising file-submission and naming conventions, automated validation (row counts, data-type and column-consistency checks), and discrepancy logging that routes issues back to client and internal teams; replaced a manual upload process with automated email-based ingestion (Microsoft Graph API to a remote SQL Server batch script), cutting weekly processing from ~4-5 hours to a ~75-minute run.
- Built and maintained 20+ automated ETL pipelines (Python, SQL, Azure) processing millions of transactions monthly and 16+ Tableau dashboards enabling spend tracking, KPI monitoring, and supplier negotiation decisions for enterprise clients.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline (Databricks, PySpark) improving detection accuracy 20%; engineered scalable transformation layers for cloud procurement data.
- Deployed a Power BI analytics solution adopted by senior Department of Transport stakeholders; presented model architecture and outcomes to cross-functional non-technical audiences.

Research Data Analyst, Terminal Ballistics Research Laboratory (TBRL), DRDO

Jan 2021 to Jul 2021

- Analysed blast-wave sensor data using statistical models (MATLAB), improving prediction accuracy 20% and contributing to 10% improvement in defence equipment design efficiency.
- Compiled findings into a 70-page structured report, translating quantitative outputs into clear, documented recommendations for defence engineering teams.

PROJECTS

MyFacit: Analytics & Decision-Support SaaS (Data Integration, Forecasting, BI, AI Automation)

- Designed and built a centralised analytics platform for a small hospitality business, integrating 8+ data sources and AI-driven workflows to automate data processing, reporting, and insight generation.
- Added short-term forecasting for operational planning and benchmarked performance against external indicators; surfaced actionable insights for pricing, product mix, and operational efficiency. Now in pilot with a local small business and in active day-to-day use, with weekly meetings to fix bugs and assess the industry-specific implications of embedding it in a real setting.
- Built the data governance layer: standardised definitions and data-quality checks that reconcile 8+ disparate sources into one consistent, trustworthy model for decision-making.

Job Hunt OS: Agentic AI Career System (Python, Claude Code, Multi-Agent)

- Designed and built a multi-agent AI system (orchestrator + 3 specialist agents) running natively on Claude Code: automates resume tailoring to JD, builds ATS-friendly PDFs, drafts cover letters, runs two-track networking (reactive applications + proactive relationship CRM), and syncs calendar reminders via MCP, with no paid APIs.
- Demonstrates production agentic system design: prompt engineering, multi-agent orchestration, file-based data modelling, scheduled cloud routines, and a human review gate at each output stage.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Coursework: Statistics I/II, Machine Learning, Applied Forecasting, High Dimensional Analysis, Communicating with Data

Bachelor of Engineering, Thapar University | May 2017 to May 2021