

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Data scientist who turns messy commercial data into decisions retail and category stakeholders can act on. I owned a pricing and cost analytics engagement end to end, presenting findings directly to a Chief Procurement Officer and category managers to help surface roughly A\$2M in savings, and built a demand forecasting model and a large-scale spend classifier (~A\$12B) that moved manual review to a repeatable pipeline. I work in Python and SQL, hold a Master of Data Science, and I am comfortable owning a problem from an ambiguous brief through to a stakeholder-ready recommendation.

KEY SKILLS

Statistics & ML: classification (logistic regression, TF-IDF, embeddings), Prophet forecasting, feature engineering, cross-validation, GridSearchCV, class-imbalance handling; familiar with gradient boosting (GBM)

Programming & Data: Python (pandas, numpy, scikit-learn), SQL, data pipelines, statistical analysis

Stakeholder & Delivery: problem framing, client and executive presentations, translating analytics into plain-language recommendations, end-to-end engagement ownership

AI & Automation: LLM workflows (Claude), agentic and multi-agent orchestration, output verification, prompt design

Cloud & BI: Azure (work), AWS (personal), Tableau, Power BI

EXPERIENCE

Data Scientist, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

A data and analytics consultancy serving enterprise clients across Australia and New Zealand. I own client engagements end to end: framing the problem, building the model, and presenting to commercial stakeholders. Selected engagements:

Pricing and cost analytics, multi-venue restaurant group (AU/NZ)

- Problem: the client's supply costs were rising but they had no visibility into where money was leaking across products, venues and suppliers.
- Approach: sole analyst on the account; built a roughly 3,000-item product catalogue and a pricing-decision dashboard (Python, SQL) tracking price movement and cost leakage down to product, venue and supplier level; presented findings and recommendations directly to the Chief Procurement Officer and category managers.
- Result: surfaced close to 30% savings, roughly A\$2M in the client's largest category, used directly in supplier negotiations.

Demand forecasting, fresh-produce and agricultural client

- Problem: the client needed to plan raw-material and chemical inventory three months ahead, exposed to shipping-cost swings.
- Approach: built a Prophet time-series model with external regressors (sea-freight trends, input prices) and delivered it through a planning dashboard.
- Result: forecasts within a 12.5 to 14% error margin, comfortably inside the client's tolerance, giving a three-month forward view for ordering decisions.

ML spend classification, enterprise procurement client (~A\$12B spend)

- Problem: only around 60 to 65% of a client's five-year, ~A\$12B spend book was reliably categorised; closing the gap took an account manager over a month per cycle.
- Approach: led the firm's first ML classifier for this problem, a one-vs-all logistic regression with features engineered from spend value, line count and transaction text (TF-IDF), tuned with cross-validation and explicit class-imbalance handling; kept the high-value head on manual review and modelled the long tail.
- Result: cut the cycle from over a month to a single day's run; documented and rolled the pipeline out across other clients.

Agentic client reporting with an independent QA agent (internal)

- Problem: analysts each spent up to 30 minutes per client, per cycle, hand-assembling progress reports across roughly 25 to 30 accounts.
- Approach: built an agentic workflow (Claude, Python) that drafts each report from analytics outputs, plus a second, independent verification agent that re-checks the key numbers before a human signs off.
- Result: rolled out company-wide, saving roughly 12 to 15 hours a week, with a trust check on every AI output before it reaches a client.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline (Databricks, PySpark) on government procurement data, improving detection accuracy by around 20%.
- Delivered findings via a Power BI solution adopted by senior Department of Transport stakeholders.

PROJECTS

Job Hunt OS: multi-agent agentic AI system (Claude Code, Python)

- A file-based multi-agent system: an orchestrator that delegates to specialist agents (resume tailoring, outreach, coaching), with a human review gate. Hands-on practice with agentic design, prompt engineering and orchestration.

Cloud Data Platform with Terraform (personal build)

- Infrastructure as code on AWS (S3, EC2, IAM) with Terraform, an automated ingestion pipeline into Snowflake, and a Power BI layer, modular across dev/staging/prod.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Statistics I/II, Machine Learning, Applied Forecasting, Supply Chain Management, High Dimensional Analysis, Communicating with Data.

Bachelor of Engineering, Thapar University | May 2017 to May 2021