

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

Hiring Team, RevenueCat

Dear Hiring Team,

I'm writing about the Data Analyst role at RevenueCat. Most of my work has been sitting inside a business team, not next to it: taking a question a stakeholder can't quite articulate yet, figuring out what they actually need to decide, and coming back with a dataset or a dashboard they use without having to ask me to explain it again. That's the job as you've described it, and it's the part of the work I'm best at and enjoy most.

The clearest proof of that is a procurement engagement I ran for a multi-venue restaurant group across Australia and New Zealand. The brief when I started was thin: the client believed they were overspending on supplier costs but had no idea where. I owned the account on my own, structured a product catalogue from scratch to make the spend analysable, then built a pricing dashboard that traces cost leakage down to product, venue and supplier. I ran the client relationship directly, presenting findings every two weeks to their Chief Procurement Officer and category managers, and the work helped them find close to A\$2M in savings in their largest category over the year, savings they used directly at the negotiating table.

The other half of what I'd bring is how I think about AI agents, because I already build with them daily and don't trust them by default. At my firm, I built an agentic workflow that drafts client progress reports from our analytics outputs, then added a second, independent QA agent whose only job is to re-derive the same numbers from a clean session, with no memory of how the first agent got there, before a human signs off and it goes to a client. That habit, treating an agent's answer as a claim to verify rather than a fact to forward, is what's let me actually rely on these tools in client-facing work instead of just experimenting with them, and it's the same instinct I'd bring to making your agentic tooling something people trust.

RevenueCat's model, remote-first across 25 or more countries, with the analytics work living close to the business rather than in a request queue, is exactly the setup I do my best work in. I don't come from subscriptions, but the underlying discipline, defining a metric precisely, explaining what a number does not say, and rebuilding trust in a number after something changes upstream, is the same discipline I've applied to procurement spend and supplier costs for the last two years. I'd welcome the chance to talk about how that transfers, and about where it doesn't yet, honestly.

Thank you for considering my application. I'd be glad to talk further at your convenience.

Regards,

Dhruv Nirmal