

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Analytics partner to business teams: I take a rough question from a stakeholder and return a number they can act on. At my consultancy I own client engagements end to end, including a procurement spend analysis that started with no dataset and surfaced close to A\$2M in savings, presented to the Chief Procurement Officer. I use AI agents daily but build the verification layer around them, an independent QA agent that re-derives every number before it reaches a client. Strong SQL, dashboards teams rely on without checking with me first, and Git as part of my daily routine.

KEY SKILLS

Analytics partnering: engagement ownership, framing ambiguous questions, presenting to senior stakeholders, clear communication of caveats, method and definition documentation

SQL & warehousing: SQL, Azure-hosted warehouses (SQL Server), data modelling, high-volume data (millions of rows, weekly refresh)

Dashboards & BI: Tableau (16+ maintained), Power BI, KPI reporting

AI workflows: Claude, multi-agent orchestration, independent QA/verification agents, human-in-the-loop

Engineering: Git (branches, PRs, review), Python, ETL automation · familiar: Snowflake (personal build), Databricks/PySpark (internship)

EXPERIENCE

Data Analyst, Purchasing Index Data Analytics (Comprara Group) | Jun 2024 to Present | Melbourne

Direct analytics partner to enterprise clients across AU/NZ, owning each engagement end to end.

- **Pricing & cost analytics, multi-venue restaurant group:** owned an ambiguous brief ("we're leaking money, we don't know where", no ready dataset); built a ~3,000-item catalogue and a Tableau price-leakage dashboard in SQL and Python, ran bi-weekly meetings, and presented to the Chief Procurement Officer and category managers. Result: close to **A\$2M (~30%)** saved in their largest category, used in supplier negotiations.
- **Agentic reporting with independent verification:** built a Claude workflow that drafts client reports from our analytics outputs, plus a second **QA agent that re-derives every number from a clean session** and a human sign-off before anything is sent. Result: adopted across ~25 to 30 accounts, ~12 to 15 hours/week saved, catching answers that looked right but weren't.
- **Data pipeline & quality, largest client (5 regions, 13 sub-regions):** own the SQL and Python pipeline behind ~6M invoice rows (~500k to 600k refreshed weekly); automated Graph-API ingestion plus row/type/column validation cut weekly processing from 4 to 5 hours to **~75 minutes**. The self-service refresh app (all 26 client accounts) lives in **Git with branches, PRs and current docs**.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI) | Aug to Nov 2023 | Melbourne

Built a distributed anomaly-detection pipeline (Databricks, PySpark), delivered via Power BI; **+~20% detection accuracy**, adopted by senior Department of Transport stakeholders.

PROJECTS

- **MyFacit (personal SaaS):** analytics product for independent cafés; unifies POS, wages and 14 suppliers' invoices into one weekly margin view plus forecasting and price-creep detection. Live with a design-partner café.
- **Cloud Data Platform (personal):** Terraform/AWS infrastructure as code with an ingestion pipeline into **Snowflake**, multi-environment (dev/staging/prod).

EDUCATION

Master of Data Science, Monash University (2022 to 2023) · **B.E.**, Thapar University (2017 to 2021)