

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61 406 259 619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

I have built and run a product end to end, deciding what it measures and building the data layer underneath it, which is the heart of this role. That product is MyFacit, a live profitability SaaS I founded for independent hospitality venues, where I define every metric it tracks and decide what the data needs to surface for a non-technical owner. Alongside it my work is SQL-first: I write and tune complex SQL (window functions, multi-table joins, query optimisation) to query and validate large transaction datasets, and I have stood up a cloud data warehouse from nothing on Snowflake with Terraform, so deciding what should sit alongside a workhorse database and building it is work I have already done. I use AI agents daily, including an independent verification pass that re-derives numbers before anything reaches a client, and I build validation into everything I ship, a habit suited to a role built around privacy and keeping data in house.

KEY SKILLS

SQL & Query Performance: Advanced SQL (window functions, multi-table joins, query optimisation), large-dataset validation, SQL Server, Snowflake

Python & Data Pipelines: Python (pandas, numpy, scikit-learn), automated ETL pipeline design, data cleaning, workflow automation

Cloud & Data Warehousing: Snowflake, AWS (S3, IAM, EC2), Terraform (IaC), Docker, Azure

AI & Agentic Workflows: Claude Code, multi-agent orchestration, independent AI verification/QA agents, LLM-assisted development, prompt and agent design

Dashboards & Product Metrics: Power BI, Tableau, KPI and metric definition, product analytics; picking up React to pair on an internal dashboard front end

Practices: data privacy and integrity discipline, Git, documentation for distributed teams

EXPERIENCE

Data Analyst / Engineer, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia

A data and analytics consultancy analysing supplier and transaction data for enterprise clients across Australia and New Zealand.

- Write and tune complex SQL, including window functions and multi-table joins, to query and validate large transaction datasets (one client pipeline runs to roughly 6 million rows) across 5+ enterprise clients, adding automated checks that catch discrepancies before they reach a report.
- Design and manage agentic AI workflows in production, including a multi-agent QA system (5 agents plus an orchestrator) that independently re-derives and verifies analytics outputs before they reach a client, using AI agents to strengthen delivery rather than replace the judgment behind it.
- Built and maintain 20+ automated ETL pipelines in Python and SQL, processing millions of procurement transactions monthly, deployed on Azure-hosted infrastructure with scheduled orchestration and no manual intervention.
- Own automated validation on the company's highest-volume client pipeline (row counts, data-type and column-consistency checks), replacing a manual upload process with automated ingestion and flagging discrepancies before they reach downstream reporting, the same integrity discipline a privacy-sensitive, internal-only product needs.

- Segment spend and transaction data by product, venue and supplier in SQL to isolate cost trends, then surface the analysis in Tableau and Power BI dashboards for procurement and finance stakeholders; this analysis helped one client identify close to A\$2M in savings in a single spend category.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built a distributed anomaly-detection pipeline in Databricks using PySpark on government procurement data, then delivered the results through a Power BI solution adopted by senior Department of Transport stakeholders; detection accuracy improved about 20%.

Research Data Analyst, Terminal Ballistics Research Laboratory (TBRL), DRDO

Jan 2021 to Jul 2021

- Analysed blast-wave sensor data with statistical models and Butterworth filtering to predict noise levels, improving prediction accuracy about 20%.

PROJECTS

MyFacit: profitability SaaS I founded (product ownership, data layer, metrics)

- Designed and built this product end to end, including deciding what it measures: unifies POS transactions, staff wages and invoices from 14 suppliers into one weekly view, then surfaces revenue, wage cost %, food cost % and gross margin alongside one recommended action for the week ahead. Also does revenue forecasting, menu engineering and supplier price-creep detection.
- Live in day-to-day use with a design-partner venue, with weekly reviews to fix issues and refine the metrics. The closest thing I have to this role: I own the product's data from the inside, deciding what it tracks and why, and building the layer that makes it useful.

Cloud Data Warehouse with Terraform (personal project): AWS, Snowflake, Terraform

- Provisioned a cloud data warehouse from scratch as infrastructure as code (S3, IAM, EC2, networking) and built an automated ingestion pipeline from external APIs into Snowflake, with modular Terraform patterns for dev, staging and prod environments. The same kind of build out a product needs once a single database stops being enough on its own.

Job Hunt OS (personal project): Claude Code, multi-agent orchestration, Python

- Designed and built a multi-agent system that runs entirely on Claude Code with no paid APIs: an orchestrator plus specialist agents for resume tailoring, outreach and career coaching, each with a built-in verification pass that re-checks numbers and attribution before anything is finalised. Run and maintain this as a system I rely on daily, not a one-off experiment, directing agents to do real work while keeping a human check on the output.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Coursework: Statistics I/II, Machine Learning, High Dimensional Analysis, Applied Forecasting.

Bachelor of Engineering, Thapar University | May 2017 to May 2021