

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

PROFESSIONAL SUMMARY

Data and analytics engineer who turns large, messy datasets into the dashboards and KPI frameworks business teams actually decide from, and owns the pipelines that keep those numbers reliable. I build and own 16+ Tableau dashboards used daily across five enterprise clients, backed by 20+ automated SQL and Python pipelines; one pricing analysis helped a client identify roughly A\$2M in savings. I work SQL-first, own problems end to end from an open question to a plain recommendation, and communicate clearly for distributed, remote teams.

KEY SKILLS

SQL & Analysis: SQL (SQL Server, Snowflake), Python, statistical analysis, ad hoc analysis, data storytelling

BI & Dashboards: Tableau, Power BI, KPI design and monitoring, dashboard development, self-service reporting

ETL & Data Pipelines: automated Python/SQL pipelines, workflow scheduling, data validation, standardised data intake

Cloud & Data Warehousing: AWS (S3, Lambda, IAM, EC2), Azure, Snowflake, Terraform (IaC), relational database design

Machine Learning (applied, as needed): scikit-learn, logistic regression, Prophet forecasting, anomaly detection

EXPERIENCE

Data & Analytics Engineer, Purchasing Index Data Analytics (Comprara Group)

Jun 2024 to Present | Melbourne, Australia (remote-capable client analytics consulting)

- Analyse large, complex supplier and transaction datasets in SQL across five enterprise clients, working independently to surface trends and cost-optimisation opportunities for procurement and finance stakeholders.
- Build and own 16+ Tableau dashboards that define and monitor core KPIs (spend, supplier performance), giving stakeholders self-service daily reporting instead of waiting on ad hoc requests.
- Build and maintain 20+ automated ETL pipelines (Python, SQL) processing millions of transactions monthly, standardising intake so the reporting stays reliable.
- Own data quality on the firm's largest account (5 global regions, ~6M invoice rows) with automated validation and consistency checks that catch discrepancies before they reach the client.
- Turn open-ended questions from category managers and executives into structured analysis and clear reports, presenting directly in bi-weekly and monthly stakeholder meetings.
- Applied ML where it adds value: a one-vs-all logistic-regression classifier (scikit-learn) that cut a client's ~A\$12B spend categorisation from ~6 weeks to a one-day run, Prophet forecasting within a 12.5-14% error margin, and anomaly detection that flagged cost leakage worth ~15% savings for one client.

Data Engineer Intern, Victorian Centre for Data Insights (VCDI)

Aug 2023 to Nov 2023 | Melbourne, Australia

- Built and presented a Power BI analytics solution adopted by senior Department of Transport stakeholders, translating model outputs into clear insights for a non-technical audience.
- Built a distributed anomaly-detection pipeline (Databricks, PySpark) and engineered scalable transformation layers to process government procurement data, improving detection accuracy by about 20%.

Research Data Analyst, Terminal Ballistics Research Laboratory (TBRL), DRDO

Jan 2021 to Jul 2021

- Applied statistical modelling and signal filtering (Butterworth filter) to blast-wave data to predict noise levels, improving prediction accuracy by about 20%, and compiled findings into a 70-page technical report.

PROJECTS

Facit (myfacit.com), profitability analytics SaaS for independent cafes (Python, multi-source data integration, forecasting, BI)

- Built a multi-tenant analytics product that unifies POS transactions, staff wages and 14 suppliers' invoices into one weekly view of profitability for independent cafe owners.
- Define and own the KPIs it tracks (revenue, wage cost %, food cost %, gross margin) and produce a weekly "pulse" with one recommended action, alongside revenue forecasting, menu engineering and supplier price-creep detection.
- Live in day-to-day use with a design-partner cafe (Neighbours Cafe, St Kilda); weekly meetings drive bug fixes and product decisions. A genuine example of owning a live product's metrics and dashboards end to end.

Cloud Data Platform (Terraform, AWS, Snowflake, Power BI)

- Provisioned cloud infrastructure as code (S3, IAM, EC2) with Terraform and built an automated ingestion pipeline from external APIs into a Snowflake data warehouse.
- Modular, multi-environment design (dev/staging/prod), reinforcing hands-on relational database and data warehouse principles on top of AWS.

EDUCATION

Master of Data Science, Monash University | Feb 2022 to Dec 2023

Coursework: Statistics I & II, Machine Learning, Communicating with Data, Applied Forecasting

Bachelor of Engineering, Thapar University | May 2017 to May 2021