

Dhruv Nirmal

Melbourne, Australia | dhruvnirmal2111@gmail.com | +61406259619 | [LinkedIn](#) | [GitHub](#)

Hiring Team, Sporty Group

Dear Hiring Team,

I am applying for the Data Scientist role at Sporty Group. I build machine learning models and statistical analyses that are meant to change a real business outcome, not just run well in a notebook, and I hold my own work to that standard whether the metric is a savings figure, a forecast error or a cut in cycle time. Sporty Group's remote-first structure also fits how I already work: independently, with clear written documentation so a distributed team stays aligned without needing to be in the same room, which matters more to me now as I plan a move to India over the coming months without wanting to lose continuity in my work.

My clearest proof of applied ML impact is a one-vs-all classification model I led for a client with roughly A\$12B of spend across five years, only 60 to 65% of which was reliably categorised to start. I built it in scikit-learn, using text features from transaction descriptions, engineered features such as spend value and line count, LogisticRegression tuned with GridSearchCV and cross-validation, and evaluated on precision, recall and F1 given the heavy class imbalance. It cut a categorisation cycle that used to take an account manager over a month down to a single day's model run, and the pipeline is now documented for reuse across other clients. Separately, a Prophet time-series forecasting model I built for a fresh-produce client, incorporating sea-freight and input-price drivers, forecast three months of inventory needs within a 12.5 to 14% error margin, and a pricing and cost-leakage analysis I owned end to end for a restaurant chain helped that client find roughly A\$2M in savings in a single year.

The other part of how I work is that I use AI tools and agents daily as part of my actual development process, not as an occasional shortcut. I designed a five-agent workflow that validates client datasets for quality issues before delivery, and a separate agentic workflow that generates structured client progress reports, both of which are now used company-wide across our analyst team. I also built a personal multi-agent system, a Job Hunt OS running natively on Claude Code, that orchestrates specialist agents to tailor resumes and cover letters and produce ATS-ready output end to end. That combination, a solid Master of Data Science foundation in statistics, ML and high-dimensional analysis, and genuine day-to-day fluency building with AI agents rather than just using them, is not something most data scientist candidates can point to.

What draws me to Sporty Group specifically is the scale and the clarity of the problem: personalisation, search and recommendation that reach a large, engaged global user base, with a stated emphasis on ML that has to prove measurable impact rather than just ship. That is the exact discipline behind the classification and forecasting work above: engineer the right features, handle imbalance, validate rigorously, and connect the result back to a number a stakeholder cares about. I would bring that same discipline to personalisation and recommendation problems, and a fully remote, product-driven environment like Sporty Group's is exactly where I want to keep building it at a larger scale.

I would welcome the chance to talk through how this experience maps onto the team's personalisation and recommendation work. Thank you for your time and consideration.

Sincerely,

Dhruv Nirmal